

# Image analysis

CS/CME/BioE/Biophys/BMI 279

Oct. 24, 2019

Adrian Sanborn

Ron Dror

# Assignment 2 tips

- Your predictors will not perfectly recapitulate the structure of HRAS, because protein structure prediction is challenging.
- But you should get a fairly compact structure.
- TAs have added extra office hours next week.

# Outline

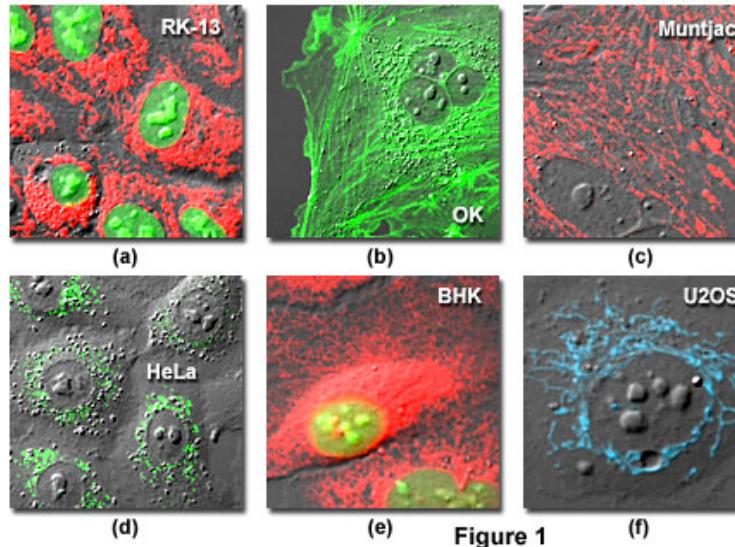
- Images in molecular and cellular biology
- Reducing image noise
  - Mean and Gaussian filters
  - Frequency domain interpretation
  - Median filter
- Sharpening images
- Image description and classification
- Practical image processing tips

# Images in molecular and cellular biology

# Most of what we know about the structure of cells come from imaging

- Light microscopy, including fluorescence microscopy

Light-microscopy uses photons to excite a fluorescent protein or dye, which will in-turn emit light of a different wavelength, allowing for visualization of protein(s), cell(s), etc labeled with the fluorescent marker. Without these markers, we would have a very difficult time interpreting images or locations within an image



<https://www.microscopyu.com/articles/livecellimaging/livecellmaintenance.html>

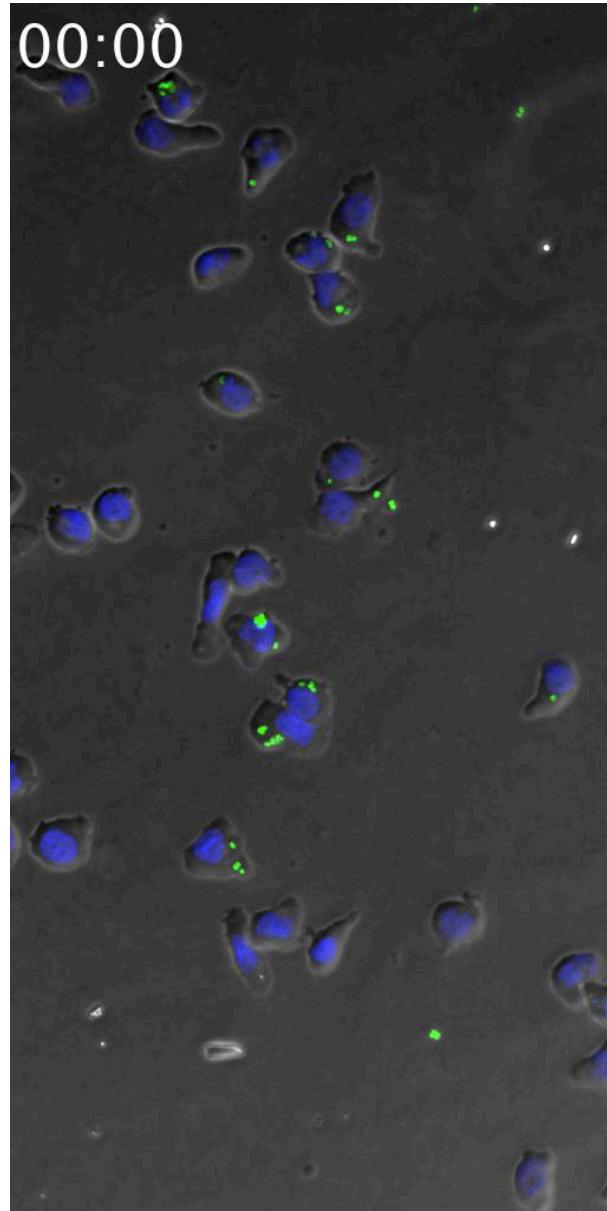
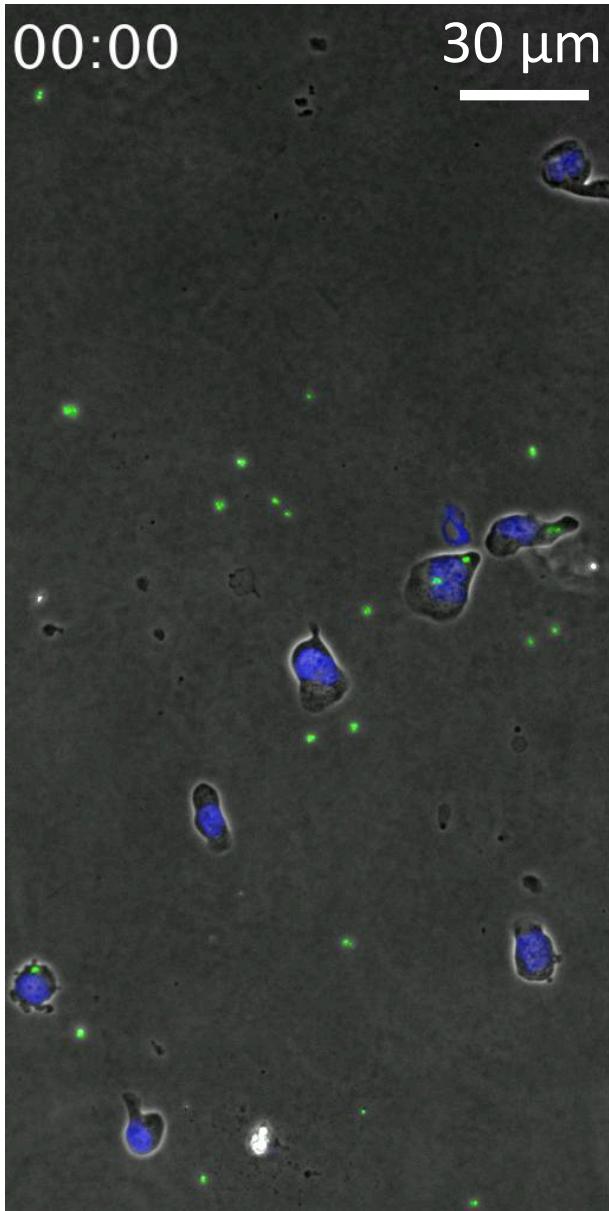
- Electron microscopy

Electron microscopy uses electrons to image a sample, providing much higher spatial resolution (on the scale of pm-nm) than light microscopy, but requiring specific conditions for the item being sampled; one such limitation is the requirement of the sample being in a vacuum, thus requiring samples to be dead.



<http://blog.library.gsu.edu/wp-content/uploads/2010/11/mtDNA.jpg>

# Imaging can capture structure over time



Human white blood cells  
eating bacteria

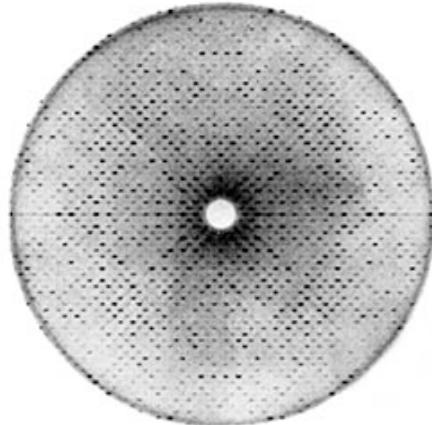
Nucleus  
*S. aureus* (GFP)

# Imaging is pervasive in structural biology

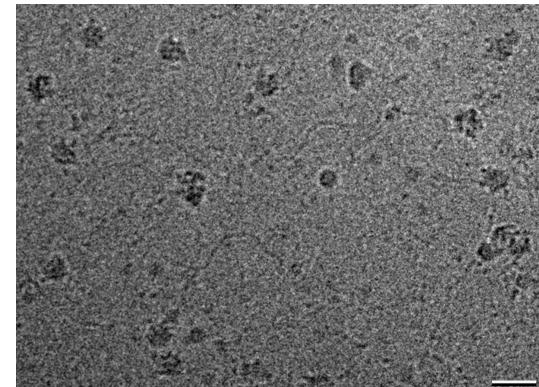
- The experimental techniques used to determine macromolecular (e.g., protein) structure also depend on imaging

The outputs of these images are less obvious, we cannot simply view the image and understand what the item being imaged looks like. In order to extract information, we have to use image processing and reconstruction tools.

X-ray crystallography



Single-particle cryo-electron microscopy



[https://askabiologist.asu.edu/sites/  
default/files/resources/articles/  
crystal\\_clear/CuZn\\_SOD\\_C2\\_x-](https://askabiologist.asu.edu/sites/default/files/resources/articles/crystal_clear/CuZn_SOD_C2_x-)

7

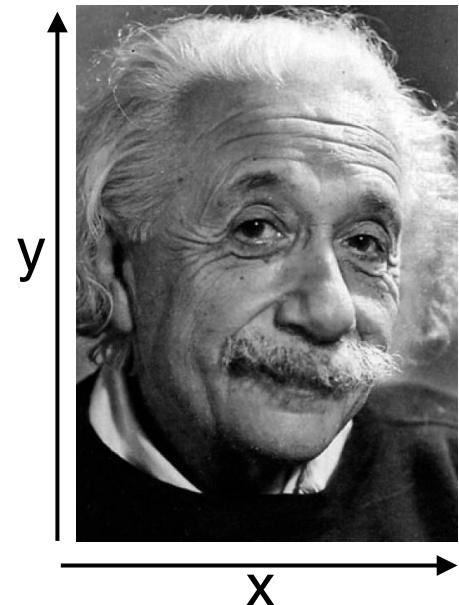
[http://debkellylab.org/?page\\_id=94](http://debkellylab.org/?page_id=94)

# Computation plays an essential role in these imaging-based techniques

- Some techniques require substantial computation before you can even see the image
- We will start with analysis of microscopy data, because it's closest to our everyday experience with images
- In fact, the basic image analysis techniques we'll cover initially also apply to normal photographs

# Representations of an image

- Recall that we can think of a grayscale image as:
  - A function of two variables ( $x$  and  $y$ )
  - A two-dimensional array of brightness values
  - A matrix (of brightness values)
- A color image can be treated as:
  - Three separate images, one for each color channel (red, green, blue)
  - A function that returns three values (red, green, blue) for each  $(x, y)$  pair

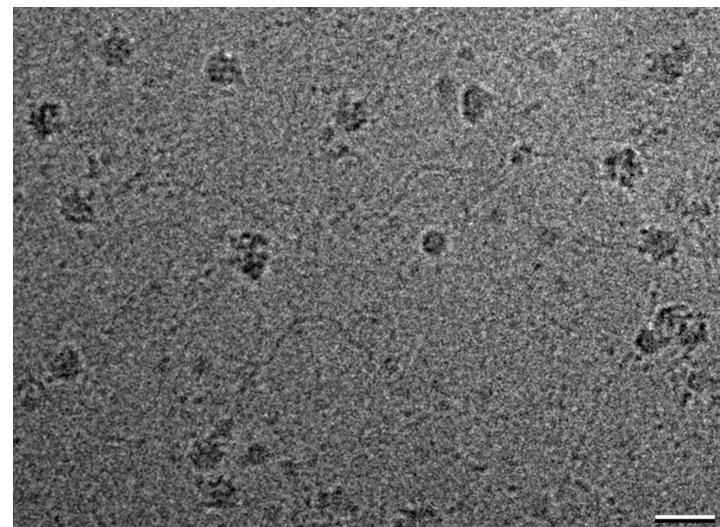


# Reducing image noise

# Experimentally determined images are always corrupted by *noise*

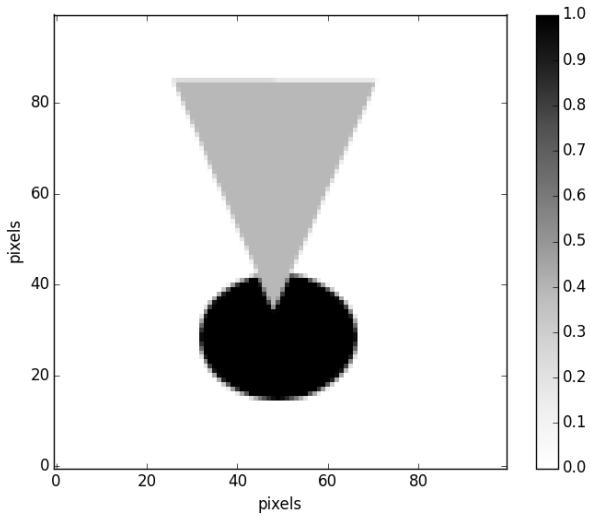
- “Noise” means any deviation from what the image would ideally look like

There is always a trade-off between signal and noise, as noise removal algorithms will also remove signal. The end-goal of noise-reduction methods is to remove as much noise as possible, while retaining as much signal as possible. Signal refers to features of an image (borders, shapes, objects, fluorescence, etc) that we are interested in analyzing.

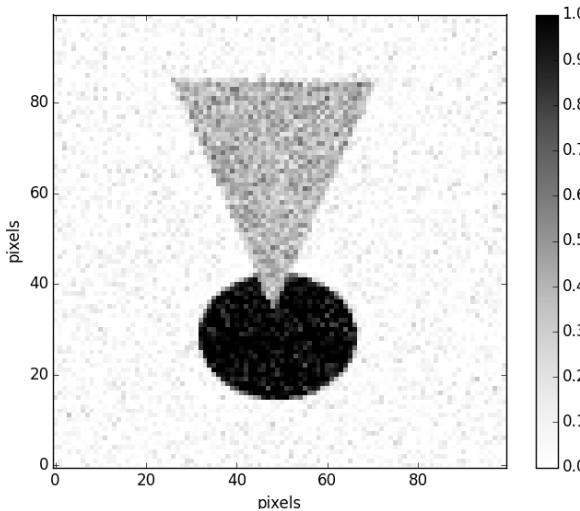


# Image noise

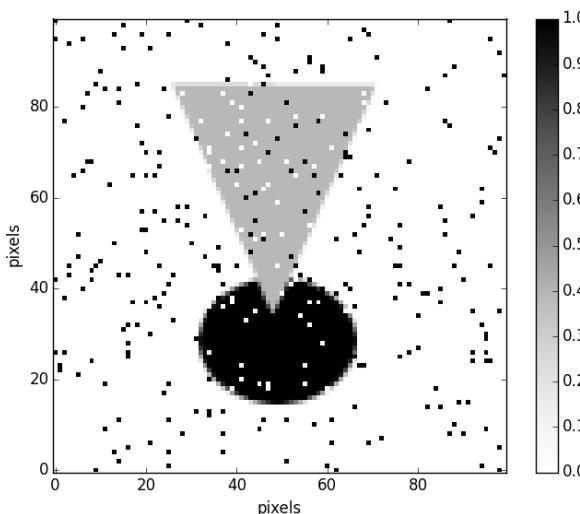
Original image



Noisy images



“Gaussian noise”: normally distributed noise added to each pixel

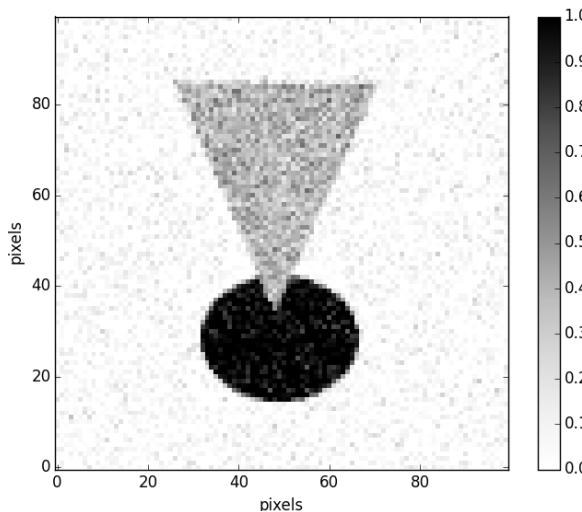


“Salt and pepper noise”: random pixels replaced by very bright or dark values

# Image noise

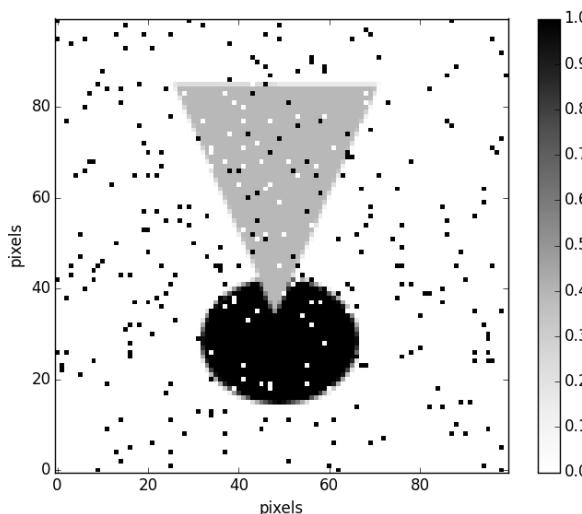
ex. Ambient light noise, aberrations, will lead to gaussian noise . This type of noise tends to be more obvious when the signal itself is relatively weak. We would say the signal-to-noise ratio (SNR) is low.

Noisy images



“Gaussian noise”: normally distributed noise added to each pixel

ex. Pixels are broken in camera or file corruption can lead to salt-and-pepper noise, in which entire pixels values are meaningless.



“Salt and peper noise”: random pixels replaced by very bright or dark values

Reducing image noise

**Mean and Gaussian filters**

# How can we reduce the noise in an image?

- The simplest way is to use a “mean filter”
  - Replace each pixel by the average of a set of pixels surrounding it
  - For example, a 3x3 mean filter replaces each pixel with the average of a 3x3 square of pixels
  - This is equivalent to convolving the image with a 3x3 matrix:  
*We could also say that we are convolving our image by a 3x3 matrix*

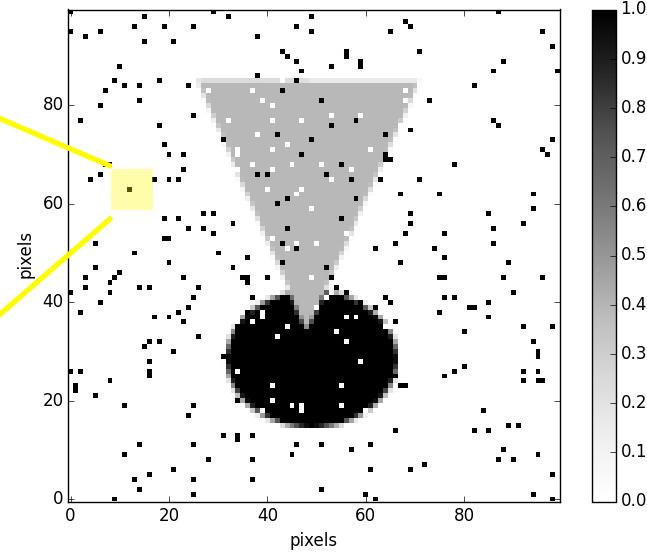
$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Why 1/9? *It's best to use an odd-sized matrix so that there is a center pixel. Then, the pixel you are examining in the original image would be at the exact center of the matrix.*  
Values should sum to 1, so that overall brightness of the image remains constant

# Mean filter

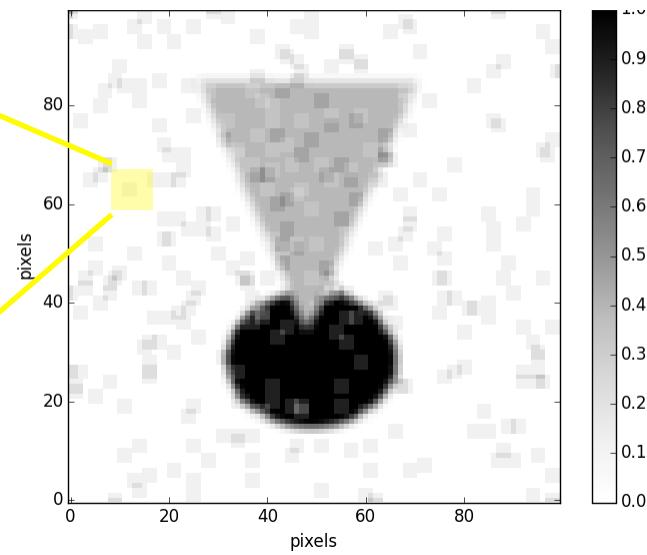
Original  
images

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



Result  
of 3x3  
mean  
filter

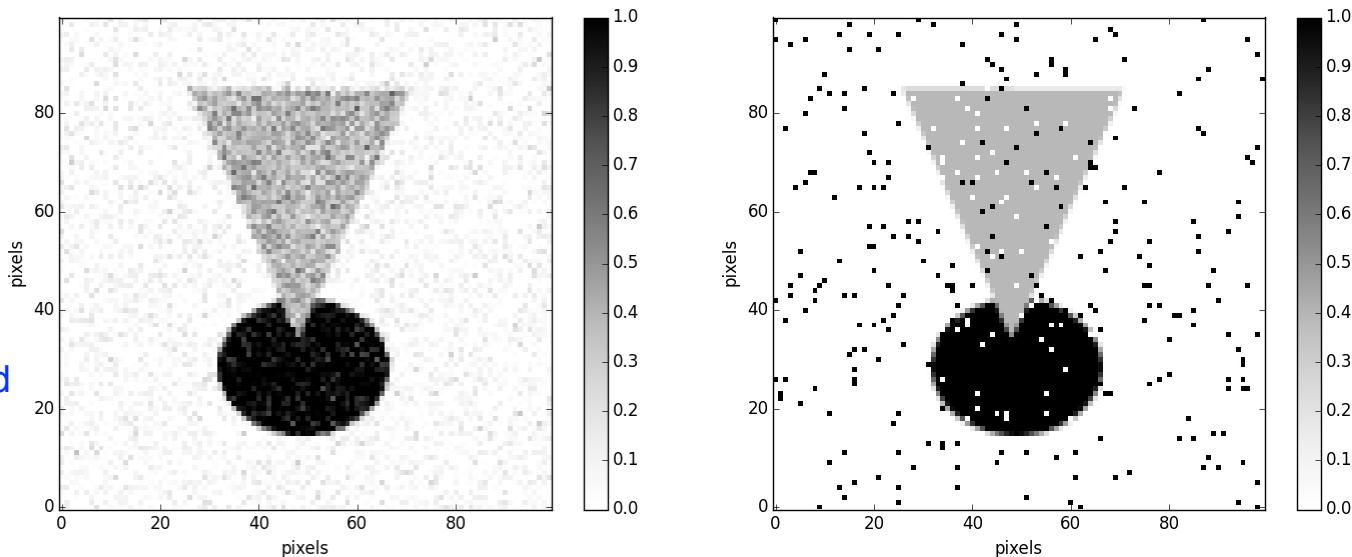
$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & \frac{1}{9} & \frac{1}{9} & \frac{1}{9} & 0 & 0 \\ 0 & 0 & \frac{1}{9} & \frac{1}{9} & \frac{1}{9} & 0 & 0 \\ 0 & 0 & \frac{1}{9} & \frac{1}{9} & \frac{1}{9} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



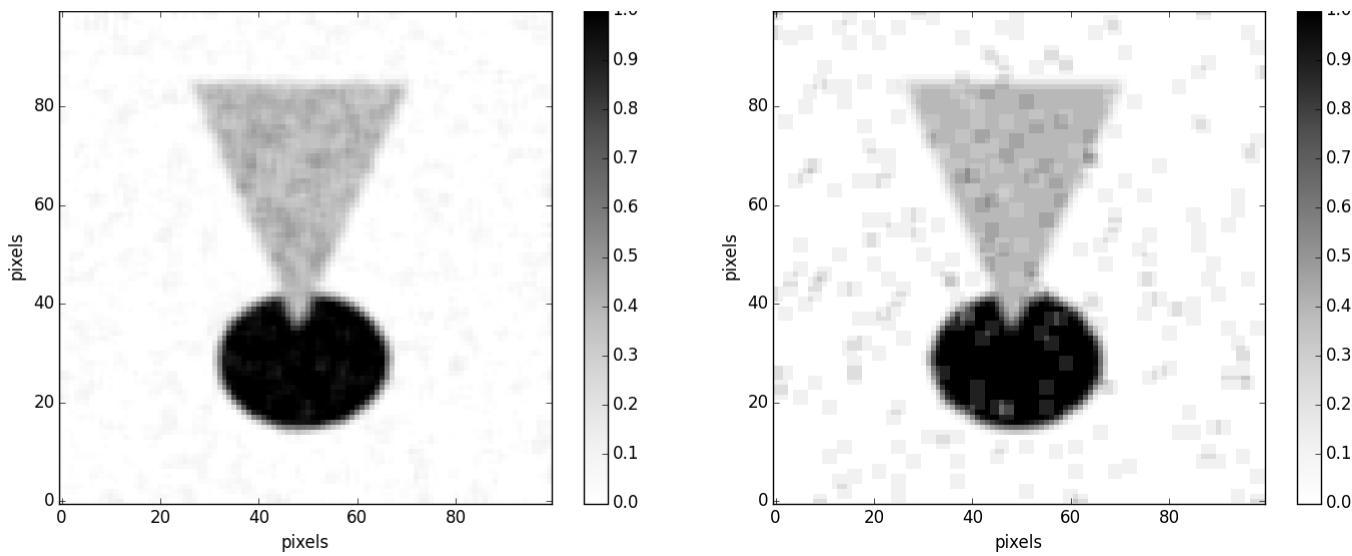
# Mean filter

Original images

Trade-off between  
Noise-reduction and  
blurring of signal



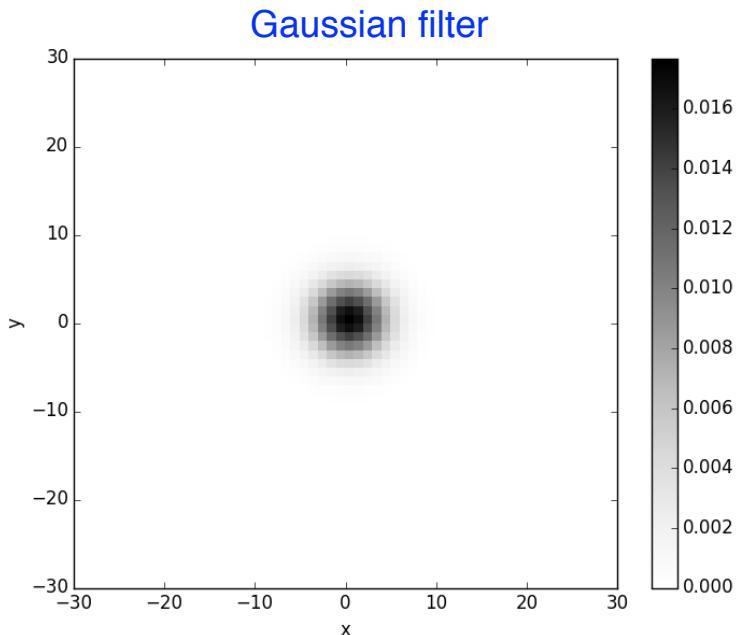
Result  
of 3x3  
mean  
filter



A larger filter (e.g., 5x5, 7x7) would further reduce noise, but would blur the image more

# A better choice: use a smoother filter

- We can achieve a better tradeoff between noise reduction and distortion of the noise-free image by convolving the image with a smoother function
- One common choice is a (two-dimensional) Gaussian
- Rather than choosing exact size, choose the standard deviation



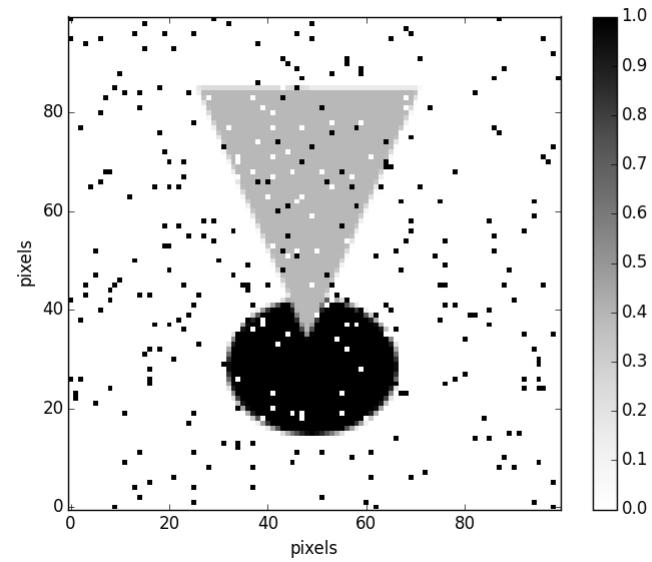
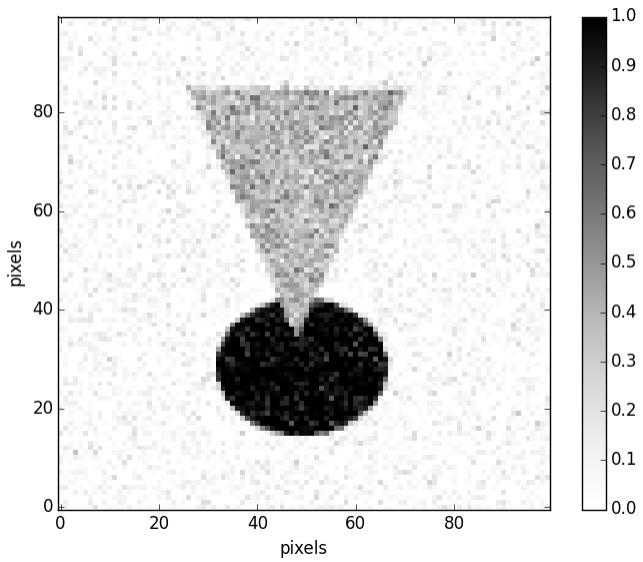
Weights sum to 1, so that the output image has the same as before.

$$\begin{bmatrix} 0.003 & 0.013 & 0.022 & 0.013 & 0.003 \\ 0.013 & 0.059 & 0.097 & 0.059 & 0.013 \\ 0.022 & 0.097 & 0.159 & 0.097 & 0.022 \\ 0.013 & 0.059 & 0.097 & 0.059 & 0.013 \\ 0.003 & 0.013 & 0.022 & 0.013 & 0.003 \end{bmatrix}$$

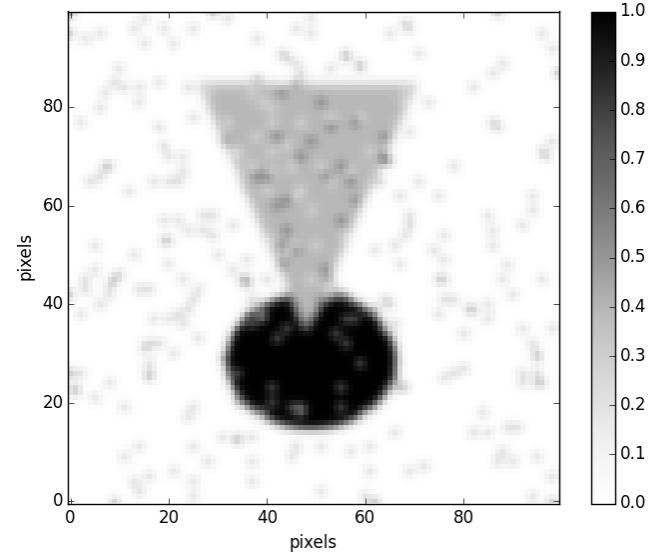
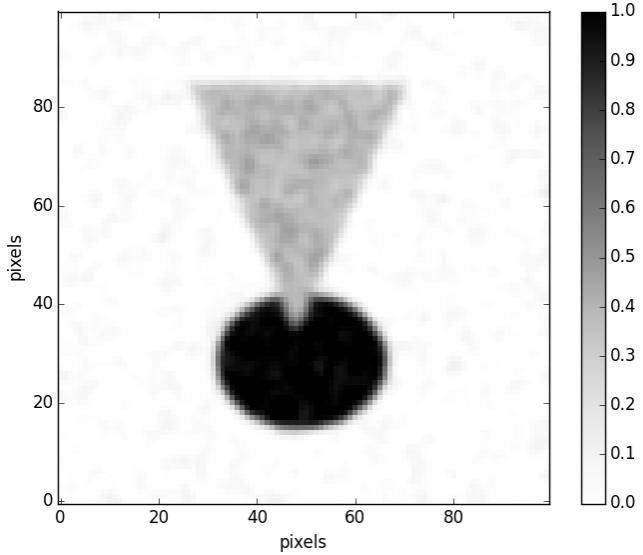
standard deviation = 1 pixel

# Gaussian filter

Original  
images

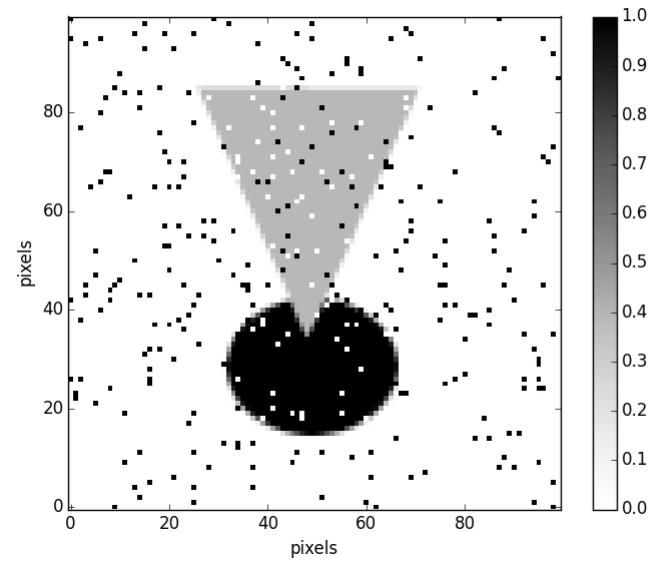
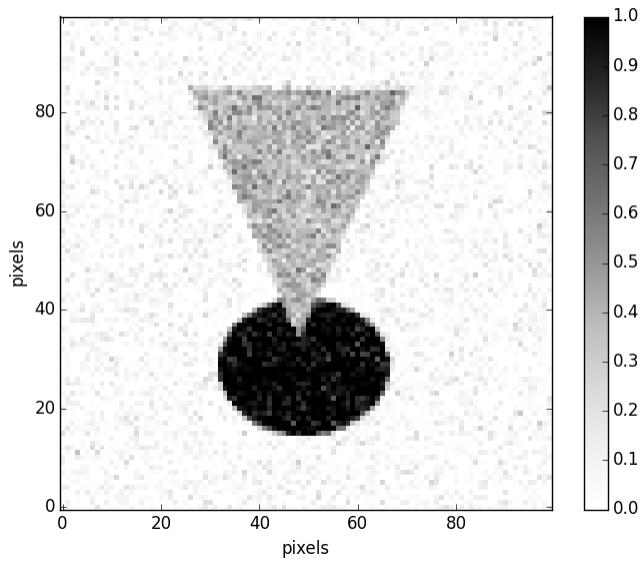


Result of  
Gaussian filter  
(standard  
deviation  
 $\sigma = 1$   
pixel)

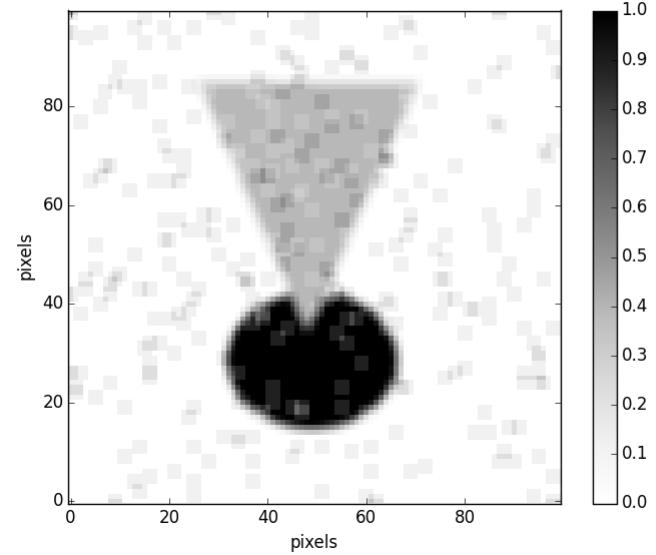
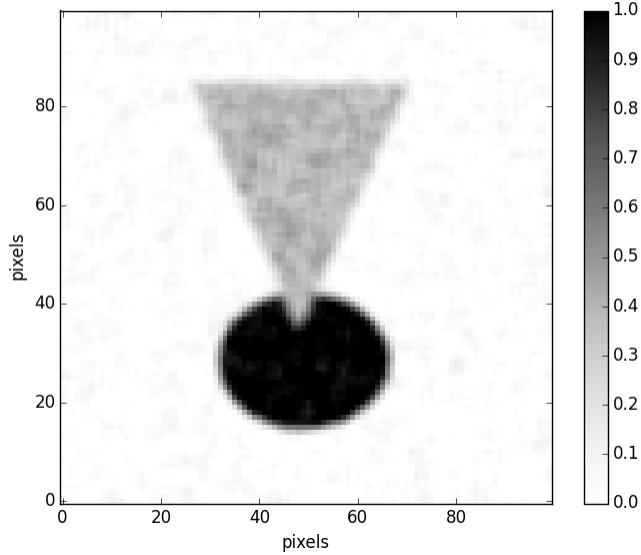


# Mean filter (for comparison)

Original  
images

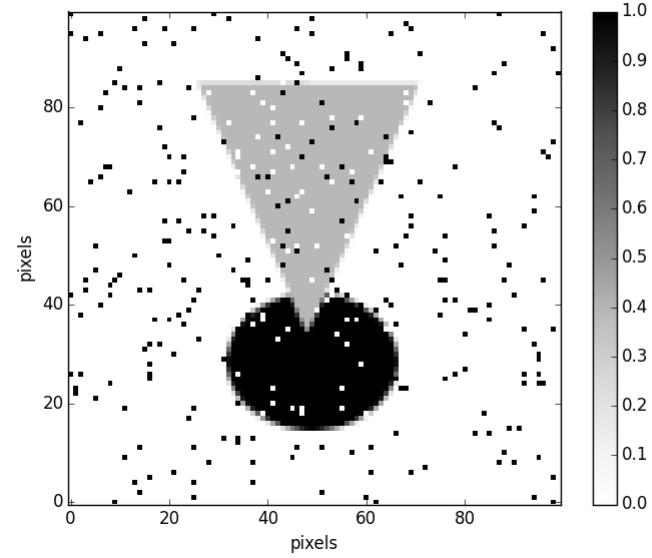
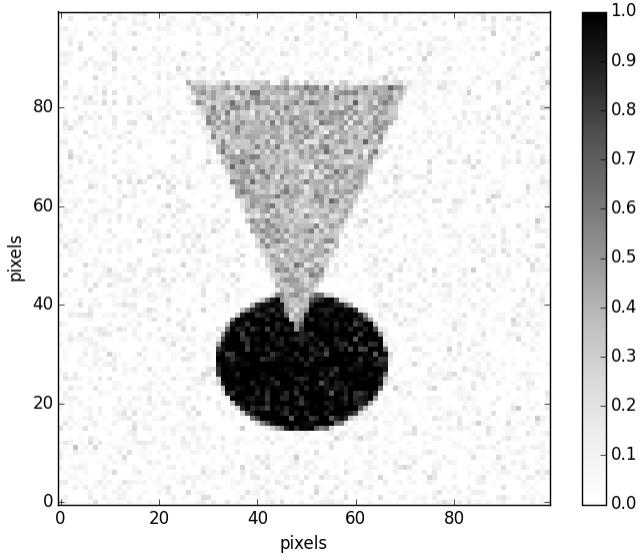


Filtered  
images

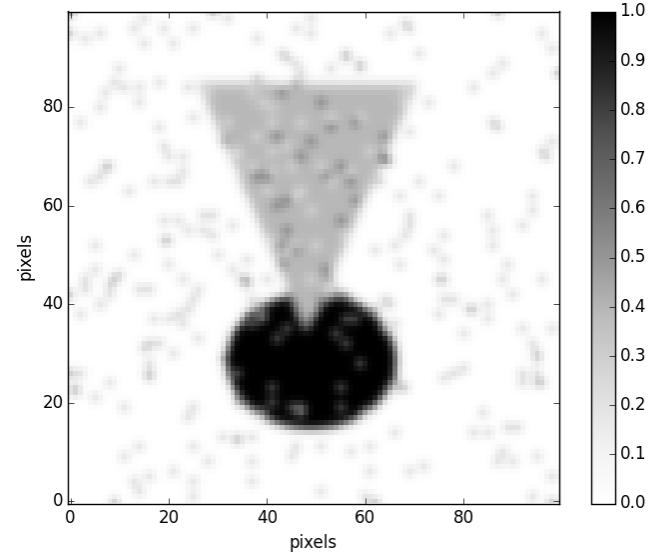
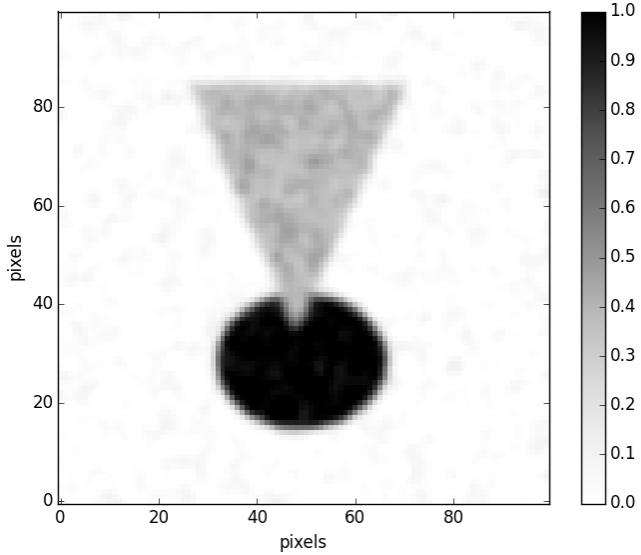


# Gaussian filter (for comparison)

Original  
images



Filtered  
images



Reducing image noise

Frequency domain interpretation

# Low-pass filtering

- Because the mean and Gaussian filters are convolutions, we can express them as multiplications in the frequency domain
- Both types of filters reduce high frequencies while preserving low frequencies. They are thus known as *low-pass filters*
- These filters work because real images have mostly low-frequency content, while noise tends to have a lot of high-frequency content

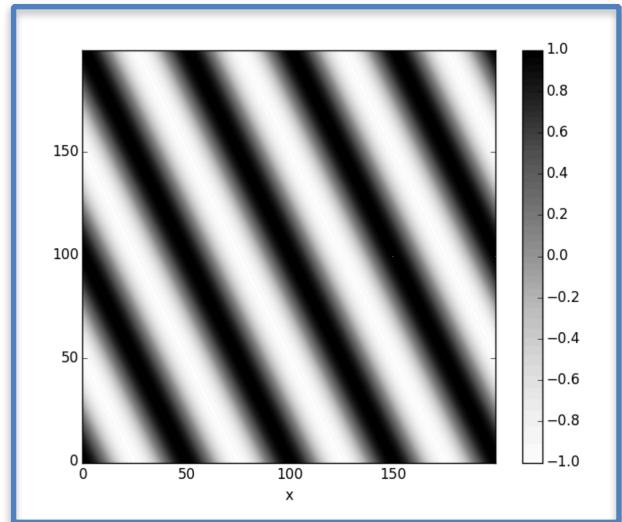
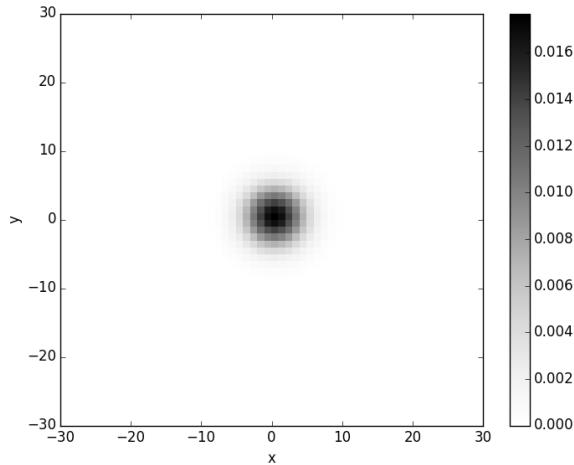
Low-frequencies pass through (are preserved) while higher frequencies do not

Note that boundaries in images are made up of higher frequencies and that is why blurring of borders/edges occurs when you apply these filters

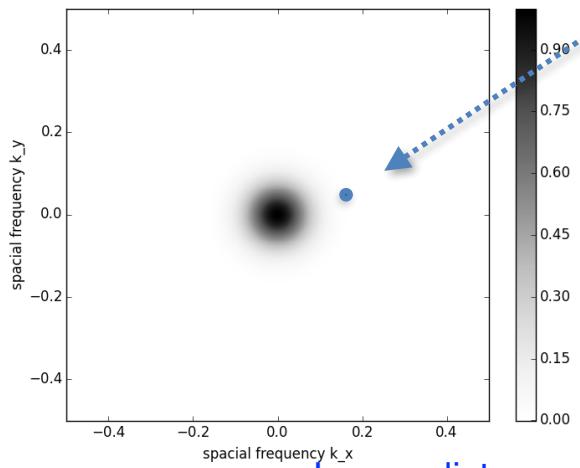
# Low-pass filtering

Filter in real domain

Gaussian filter



Magnitude profile  
in frequency  
domain  
(low frequencies  
are near center of  
plots)



Point value in the frequency  
domain describes the  
strength of the corresponding  
frequency in the real domain

Larger distance away from origin indicates higher frequencies

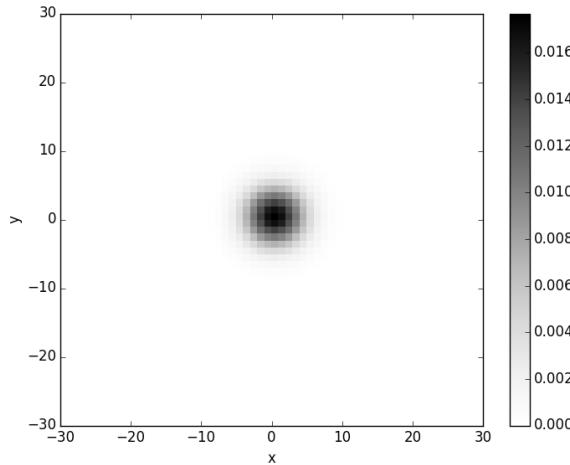
Angle of vector from origin represents direction of sine wave, or the relative x and y frequencies

Fourier transform of a gaussian (in the position domain) turns out to also be a gaussian (in the frequency domain)

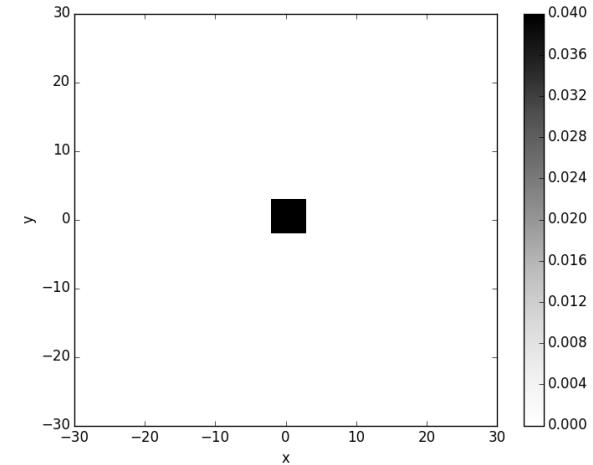
# Low-pass filtering

Filter in real domain

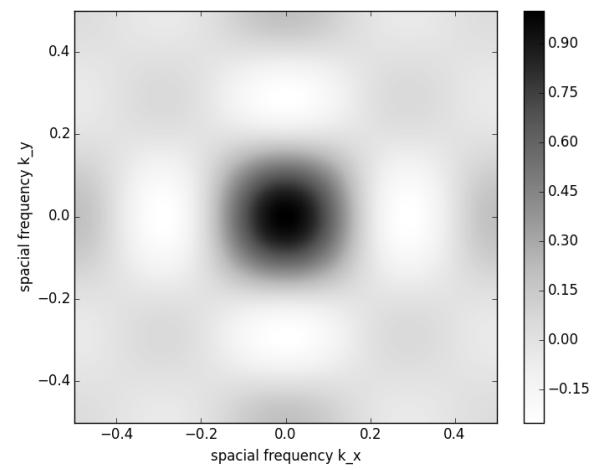
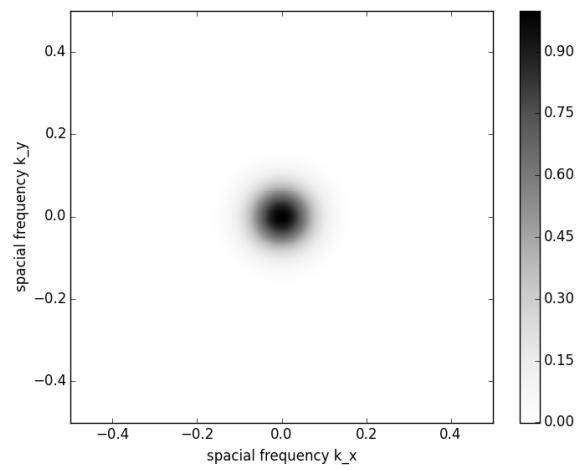
Gaussian filter



Mean filter



Magnitude profile  
in frequency  
domain  
(low frequencies  
are near center of  
plots)



The Gaussian filter eliminates high frequencies more effectively than the mean filter, making the Gaussian filter better by most measures.

# Low-pass filtering

- As a filter becomes larger (wider), its Fourier-domain representation becomes narrower
- In other words, making a mean or Gaussian filter larger will make it more low-pass (i.e, narrow the range of frequencies it passes)
  - Thus it will eliminate noise better, but blur the original image more

# Reducing image noise

## Median filter

# Median filter

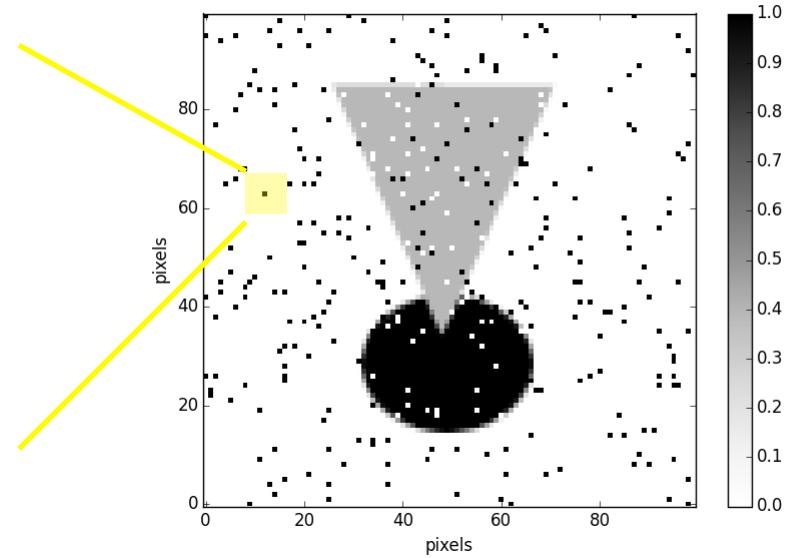
- A median filter ranks the pixels in a square surrounding the pixel of interest, and picks the middle value
- This is particularly effective at eliminating noise that corrupts only a few pixel values (e.g., salt-and-pepper noise)
- This filter is *not* a convolution

^ This is because a convolution is applied the exact same way on each region/window of the image, however, the median operations depends on what values are within that region, and thus will not be the same each time

# Median filter

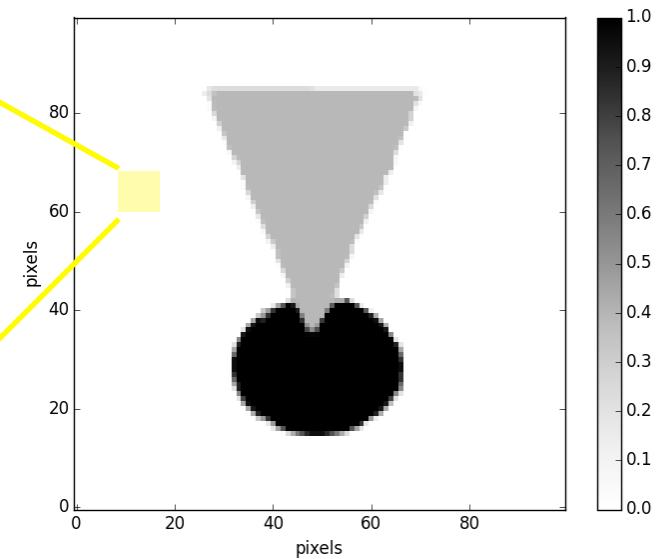
Original images

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



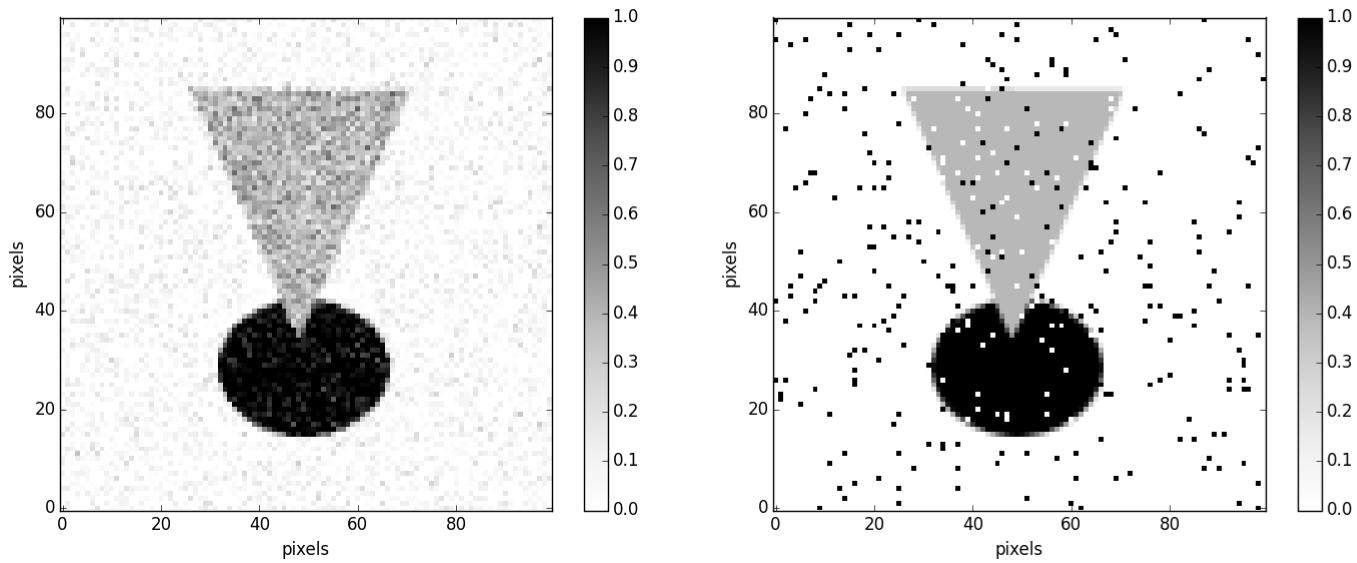
Result of  
3x3  
median  
filter

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

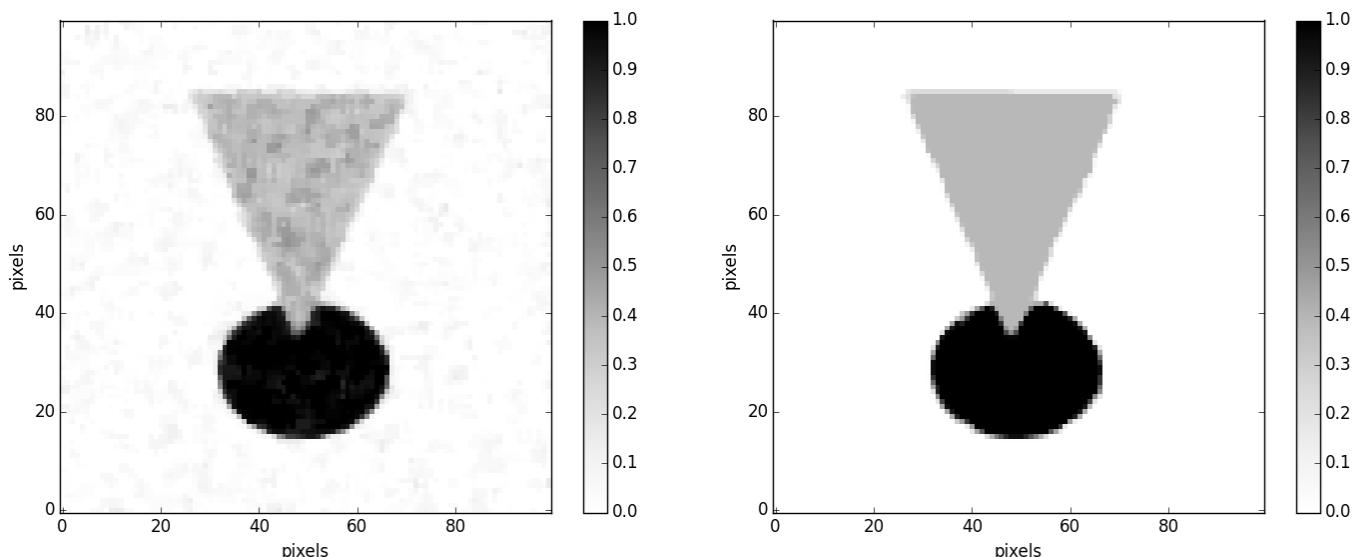


# Median filter

Original images



Result of  
3x3  
median  
filter



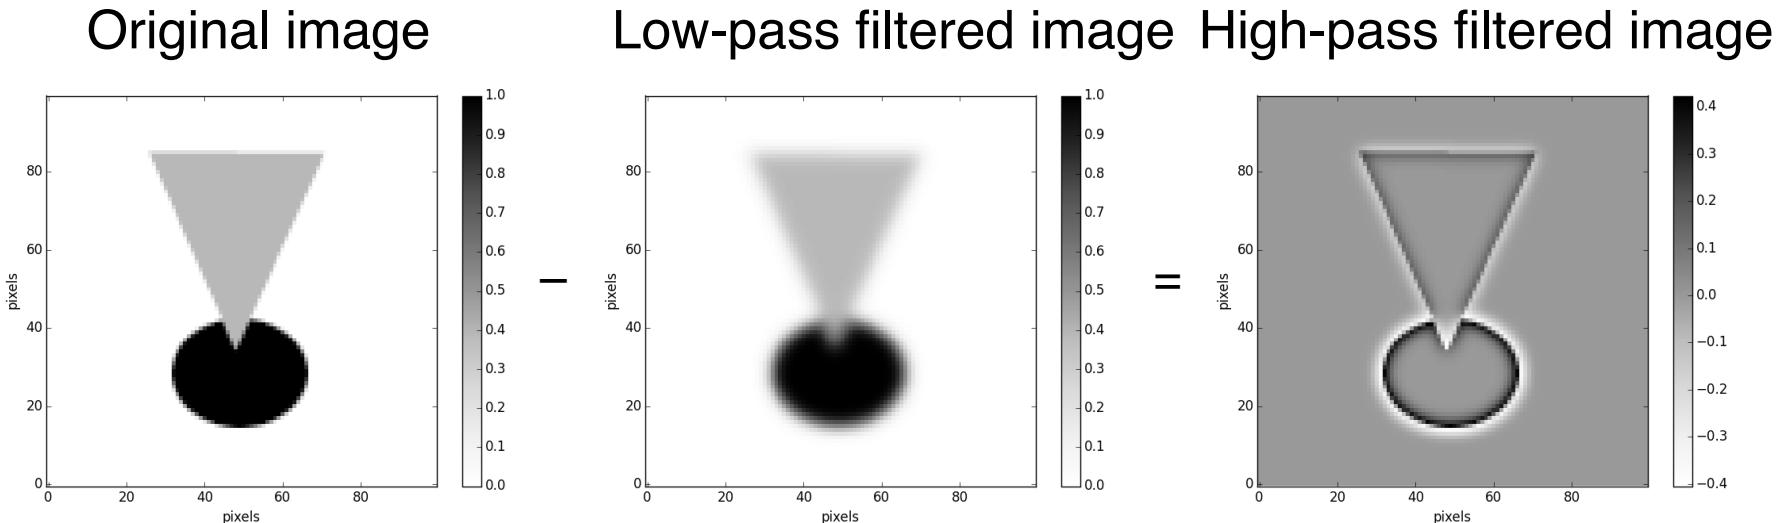
Using a larger window would further reduce noise, but would distort the image more

# Sharpening images

# High-pass filter

- A *high-pass filter* removes (or reduces) low-frequency components of an image, but not high-frequency ones
- The simplest way to create a high-pass filter is to subtract a low-pass filtered image from the original image
  - This removes the low frequencies but preserves the high ones
  - The filter matrix itself can be computed by subtracting a low-pass filter matrix (that sums to 1) from an “identity” filter matrix (all zeros except for a 1 in the central pixel)

High-pass filtering is the same as subtracting  
the low-pass image from the original



# High-pass filter

Horizontal edge  
in original image

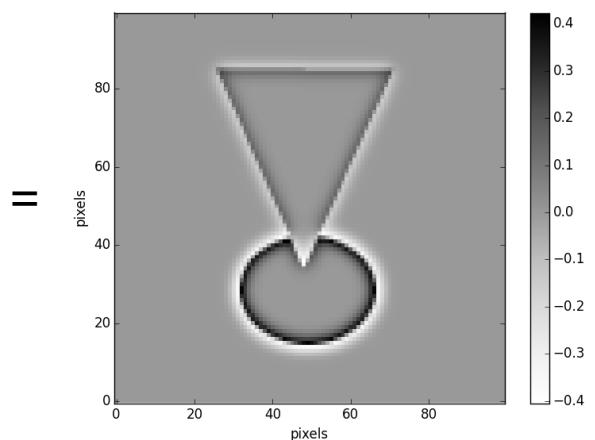
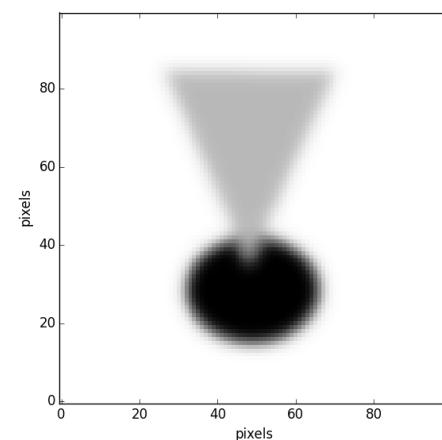
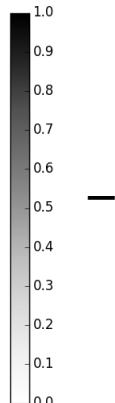
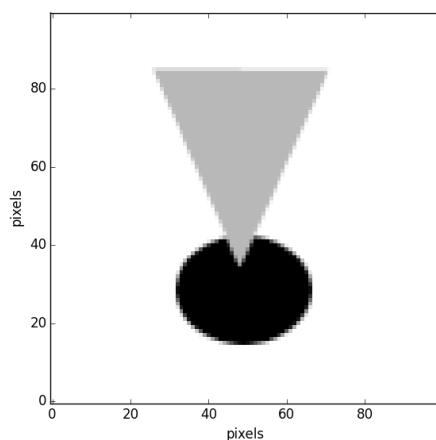
$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 9 & 9 & 9 & 9 & 9 \\ 9 & 9 & 9 & 9 & 9 \\ 9 & 9 & 9 & 9 & 9 \end{bmatrix}$$

Low-pass filtered  
( $3 \times 3$  mean filter)

$$- \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 3 & 3 & 3 & 3 & 3 \\ 6 & 6 & 6 & 6 & 6 \\ 9 & 9 & 9 & 9 & 9 \\ 9 & 9 & 9 & 9 & 9 \end{bmatrix} =$$

High-pass filtered  
(edges are emphasized)

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ -3 & -3 & -3 & -3 & -3 \\ 3 & 3 & 3 & 3 & 3 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



# How might one use a high-pass filter?

- To highlight edges in the image
- To remove any “background” brightness that varies smoothly across the image  
    ^ This kind of effect is low-frequency
- Image sharpening
  - To sharpen the image, one can add a high-pass filtered version of the image (multiplied by a fractional scaling factor) to the original image
  - This increases the high-frequency content relative to low-frequency content
  - In photography, this is called “unsharp masking”

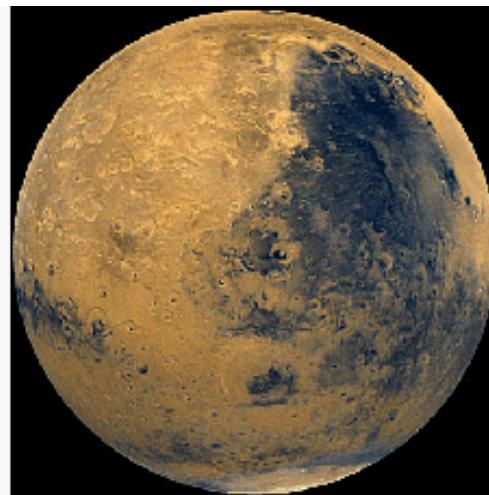


# Image sharpening — another example

Original image



Sharpened image



# Image description and classification

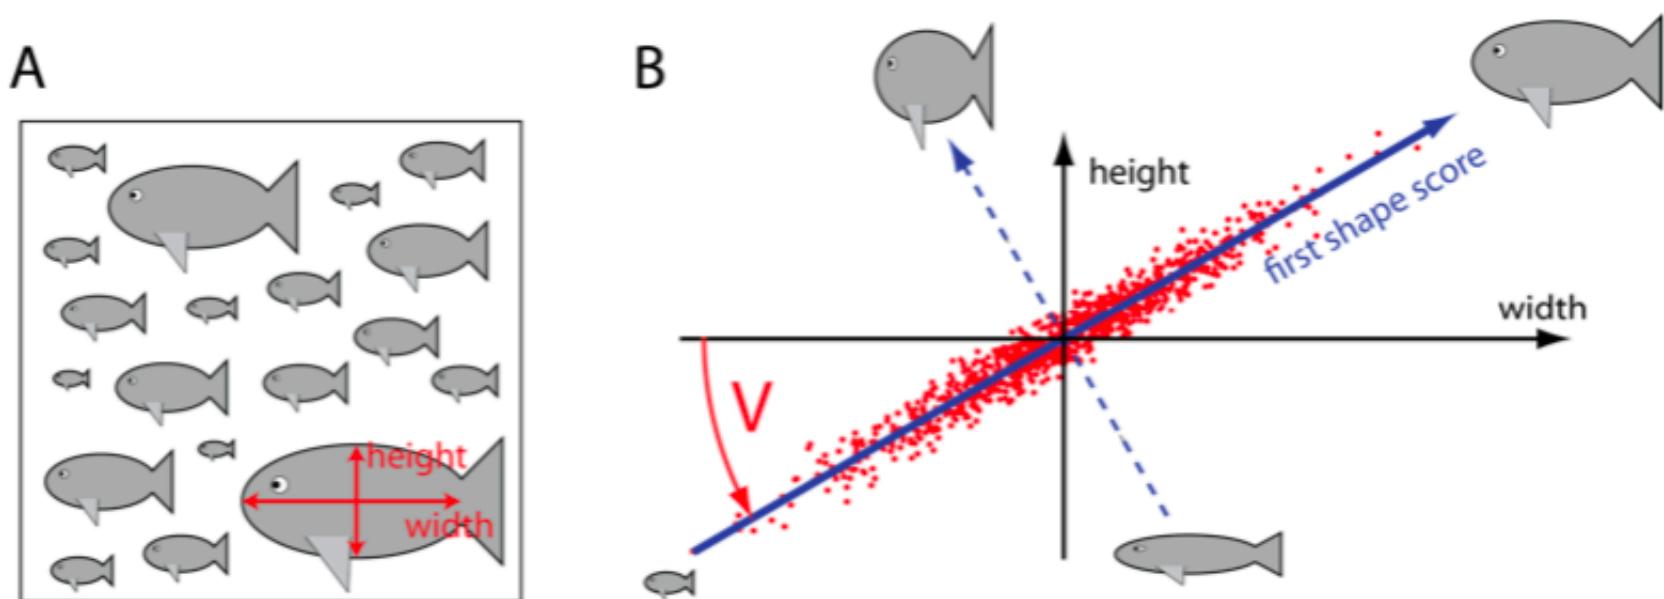
# Describing images concisely

- The space of all possible images is very large. To fully describe an  $N$ -by- $N$  pixel grayscale image, we need to specify  $N^2$  pixel values.
- We can thus think of a single image as a point in an  $N^2$ -dimensional space.
- Classifying and analyzing images becomes easier if we can describe them (even approximately) with fewer values.
- For many classes of images, we can capture most of the variation from image to image using a small number of values
  - This allows us to think of the images as points in a lower-dimensional space
  - We'll examine one common approach: Principal Components Analysis.

How can we compress the images down to a lower dimensional space?

# Principal component analysis (PCA)

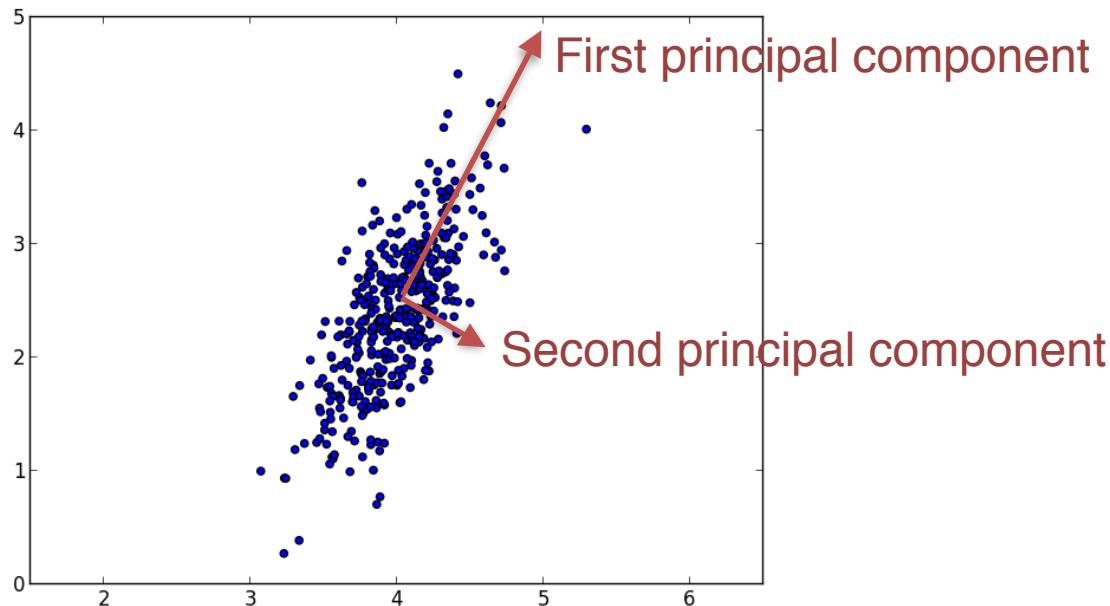
- Basic idea: given a set of points in a multi-dimensional space, we wish to find the linear subspace (line, plane, etc.) that best fits those points.



How can we define the parameters that capture the most information in the data?

# Principal component analysis (PCA)

- Basic idea: given a set of points in a multi-dimensional space, we wish to find the linear subspace (line, plane, etc.) that best fits those points.

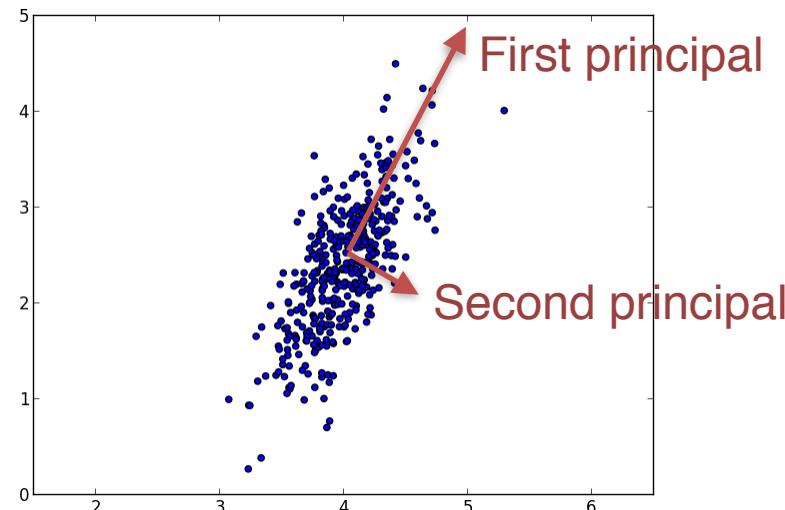


- If we want to specify a point  $\mathbf{x}$  with just one number (instead of two), we can specify the closest point to  $\mathbf{x}$  on the line described by the first principal component (i.e., project  $\mathbf{x}$  onto that line).
- In a higher dimensional space, we might specify the point lying closest to  $\mathbf{x}$  on a plane specified by the first two principal components.

# Principal component analysis (PCA)

- How do we pick the principal components?
  - First subtract off the mean value of the points, so that the points are centered around the origin.
  - The first principal component is chosen to minimize the sum squared distances of the points to the line it specifies.
  - This is equivalent to picking the line that maximizes the variance of the full set of points after projection onto that line. ie, we want to pick a line, such that there is the most "spread"/greatest range of data points on that line
  - The  $k$ th principal component is calculated the same way, but required to be orthogonal to previous principal components

Enforcing orthogonality means that the variation/information we capture in each component won't overlap

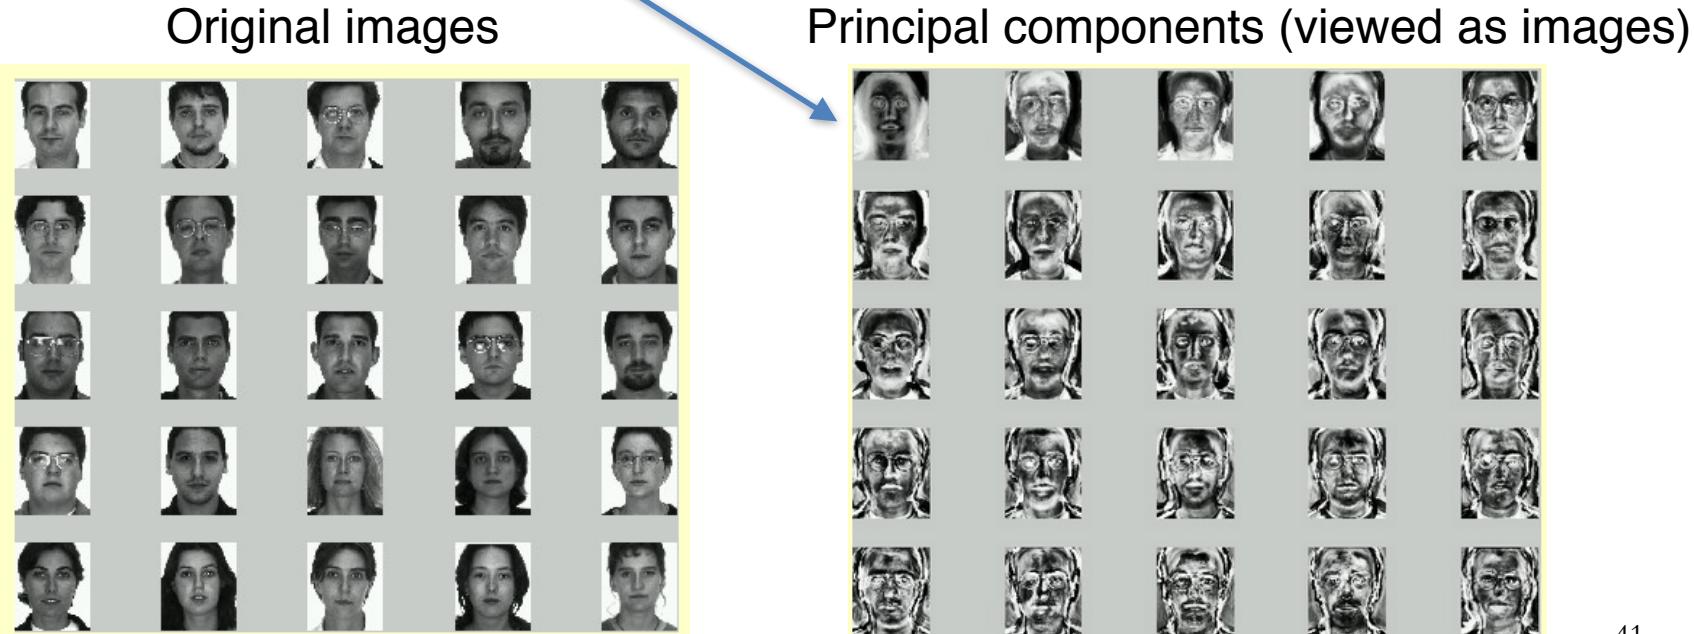


We stop finding additional components once the variation explained by the next component falls below a certain threshold. Ie we've captured most of the variance in the data using the components we've already found

# Example: face recognition

- A popular face recognition algorithm relies on PCA
  - Take a large set of face images (centered, frontal views)
  - Calculate the first few principal components
  - Approximate each new face image as a sum of these first few principal components (each multiplied by some coefficient).
  - Classify faces by comparing these coefficients to those of the original face images

First component represents signal from face and neck, but not from background



41

<http://www.pages.drexel.edu/~sis26/Eigenface%20Tutorial.htm>

We are capturing which aspects of the image are most important to recognizing and distinguishing different people

# Practical image processing tips

# Practical image processing tips

(Thanks to Leo Kesselman)

- When viewing an image, you might need to scale all the intensity values up or down so that it doesn't appear all black or all white
- You might also need to adjust the “gamma correction” to get the image to look right
  - This applies a nonlinear (but monotonically increasing) function to each pixel value
- ImageJ is a great program for looking at scientific images